राष्ट्रीय प्रौद्योगिकी संस्थान, रायपुर (छ.ग.), भारत

National Institute of Technology, Raipur (Chhattisgarh), India

Department of Information Technology

Minor Project Final Presentation — November 2026

Spatial-Guided Spectral-Channel Fusion Network for Hyperspectral Remote Sensing Image Analysis

Under the Guidance of

Dr. R.R. Janghel

Associate Professor (Supervisor)
Dept. of Information Technology,
National Institute of Technology, Raipur

Presented By

Tejas Surya Roll No. 23118098

B. Vyshnavi Roll No. 23118021

Valluri Lahitha Roll No. 23118106

B.Tech Seventh Semester, Nov 2026

National Institute of Technology, Raipur

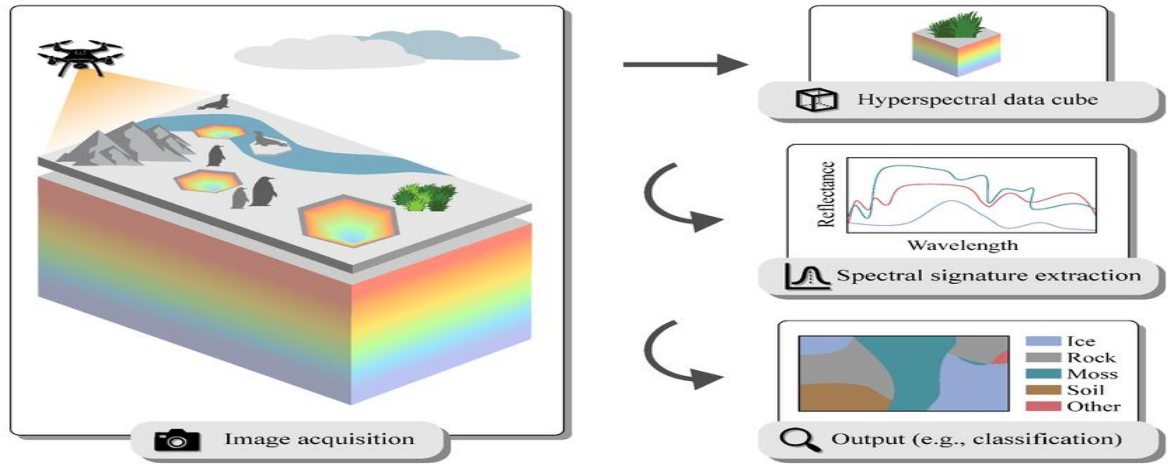
Outline

- Recap: Problem Statement & Research Objectives
- SGSCF-Net Architecture Overview
- Module Design: SBAM, CBSA, SCFG
- Hybrid Loss Function
- Experimental Setup & Benchmark Datasets
- Results: Indian Pines, Pavia University, Salinas Valley
- Classification Maps Comparison
- Computational Efficiency & Parameter Analysis
- Ablation Study & Depth Sensitivity
- Hyperparameter Grid Search (60 Runs)
- Data Scarcity Stress Testing (1%–10%)
- Cross-Dataset Summary & Key Achievements
- Future Work & Research Directions

Recap: Problem Statement & Research Objectives

Problems Identified (Mid-Semester)

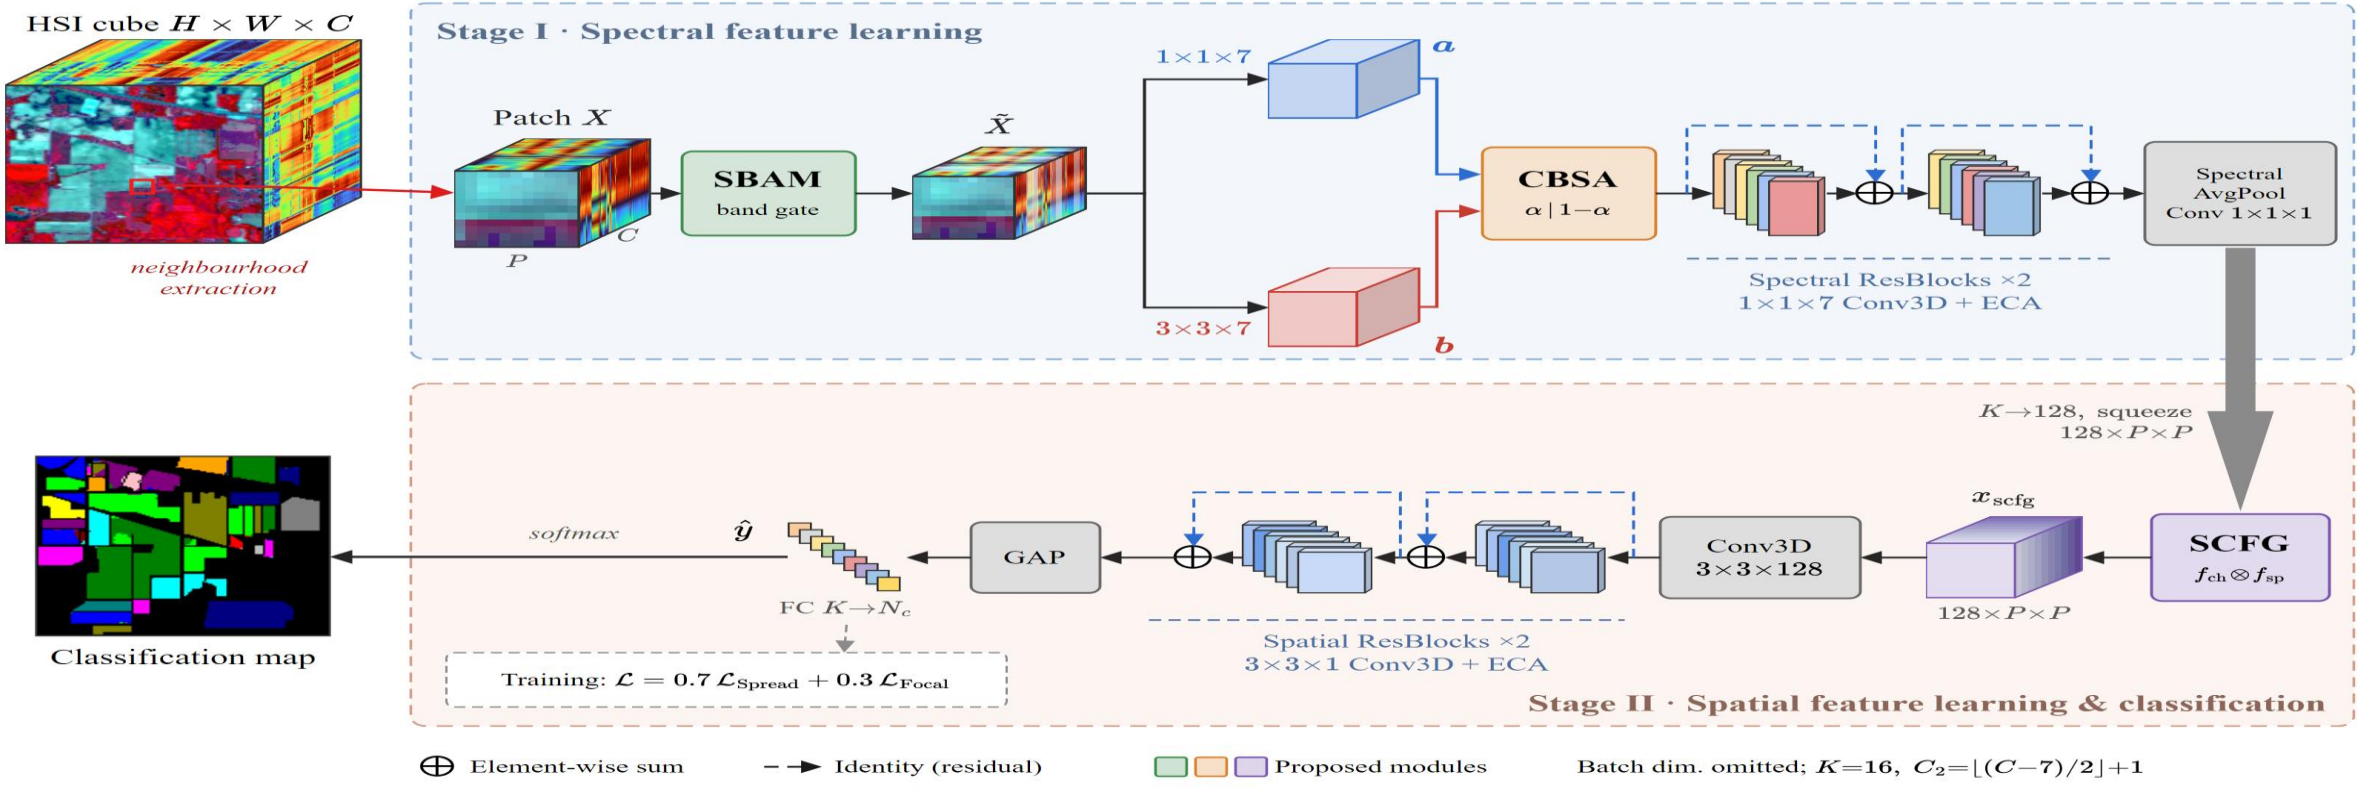
- Hughes Phenomenon: High bands + low samples = overfitting
- Memory Explosion: 3D-CNNs & Transformers exceed 85M+ parameters
- Feature Entanglement: Spatial + spectral mixing blurs class signatures
- Class Imbalance: Minority classes (Oats: 20 samples) collapse to 0%
- Unfiltered Noise: Atmospheric absorption corrupts raw band features



What We Delivered (This Presentation)

- Fully implemented & trained SGSCF-Net across 3 benchmark HSI datasets (Indian Pines, Pavia University, Salinas Valley)
- 120+ GPU training runs: 60-run patch/batch grid search + 12 data scarcity stress tests + 24 baseline benchmark runs
- Competitive or superior accuracy vs. 8 SOTA baselines while using only 308K parameters (vs. 6.4M for SSAN, ~610K for ViT)
- Complete 31-page ACM-format research paper with per-class breakdowns, classification maps, ablation studies, convergence analysis

SGSCF-Net: End-to-End Architecture Overview



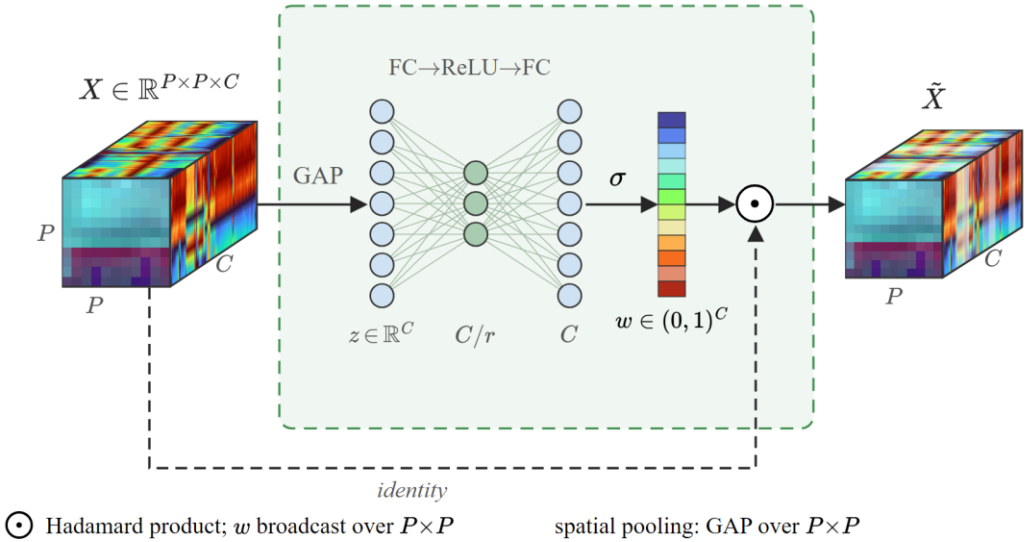
Input: X ∈ ℝ^(B×1×P×P×C) → SBAM recalibration → Dual-branch CBSA → Spectral deepening → SCFG covariance gating → Spatial refinement → ŷ ∈ ℝ^(B×C_cls)

Stage I: SBAM band gate → Stage II: Dual 3D entry + CBSA complementary gating → Spectral ResBlocks ×2 + ECA → Stage III: SCFG covariance transition → Stage IV: Spatial ResBlocks → GAP + FC | Loss: 0.7×Spread + 0.3×Focal

Module 1: Spectral Band Attention Module (SBAM)

What SBAM Does

- Operates on RAW spectral bands BEFORE any convolution
- Adaptive 3D average pooling compresses P×P spatial grid
- Squeeze MLP bottleneck (reduction r=8) learns band weights
- Sigmoid gate w ∈ (0,1)^C suppresses noisy atmospheric bands
- Hadamard product: $\tilde{X} = X \odot w$ (channel-wise reweighting)



Key Equations & Design Insight

- $z_c = (1/P^2) \sum X(b,1,u,v,c)$ [spatial pooling]
- $w = \sigma(W_2 \cdot \delta(W_1 z + b_1) + b_2)$ [MLP bottleneck]
- $\tilde{X} = X \odot E(w)$ [band filtering]
-
- Recalibrates BEFORE convolution (unlike SENet/CBAM)
- Only 10,225 params (3.32% of total network)

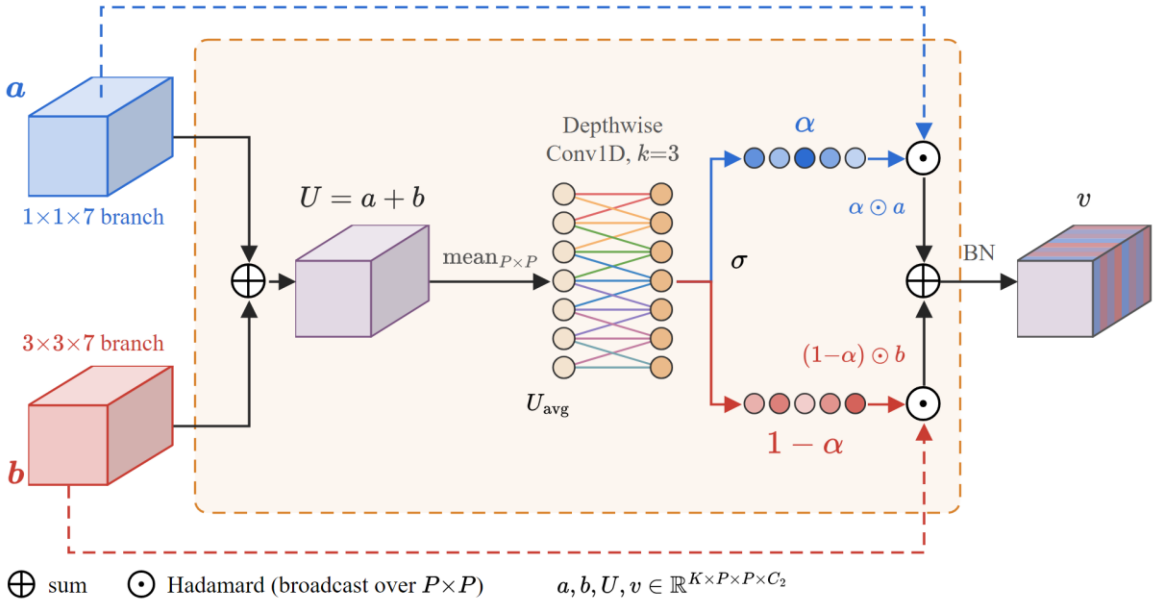
Module 2: Complementary Branch Spectral Attention (CBSA)

Dual-Branch Design

- Branch a: 1×1×7 Conv3D → pure spectral inter-band transitions
- Branch b: 3×3×7 Conv3D → joint spatial-spectral context
- Both: K=16 filters, spectral stride=2, BN + ReLU
-
- Complementary Gating:
- α = σ(Conv1D_k(MeanPool(a + b))) [k=3]
- v = BN(α ⊙ a + (1-α) ⊙ b)

Key Insight & Efficiency

- α + (1-α) ≡ 1 → ZERO redundancy between branches
- When α → 1: spectral fidelity dominates
- When α → 0: spatial context dominates
- Forces non-overlapping orthogonal representations
- Only 80 parameters (0.03%) for the gating layer
- Followed by 2 Spectral ResBlocks + ECA-3D (7,364 params)



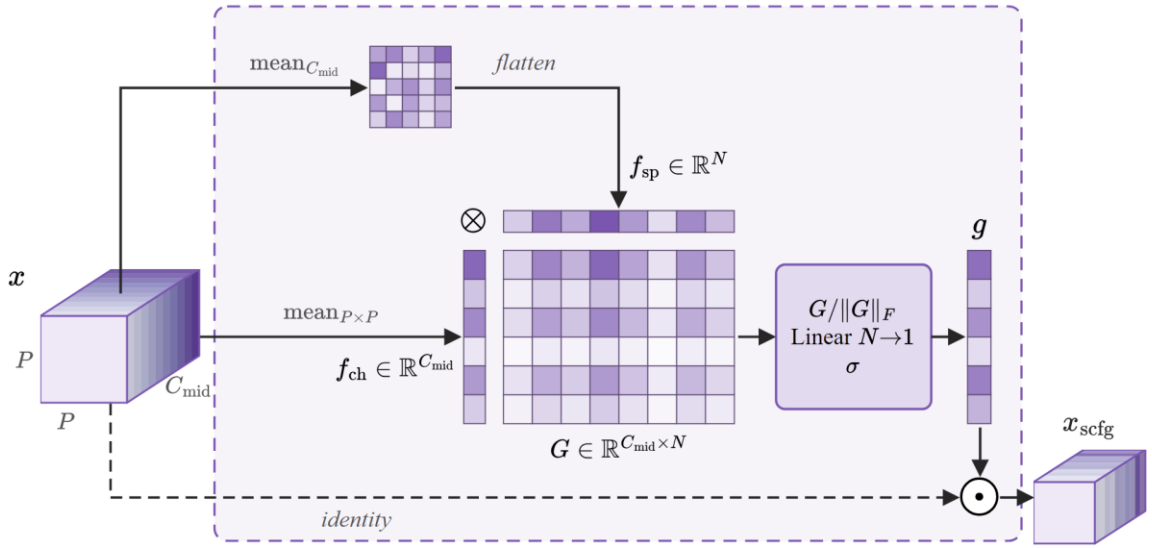
Module 3: Second-Order Covariance Fusion Gate (SCFG)

Covariance-Based Channel Calibration

- Transitions from spectral → spatial processing stage
- Adaptive 3D pool collapses spectral dim: $C_2 \rightarrow 1$
- 1×1×1 Conv3D expands: K=16 → C_mid=128
-
- Second-Order Statistics:
- f_ch ∈ ℝ^(B×C): channel descriptor (mean over spatial)
- f_sp ∈ ℝ^(B×N): spatial descriptor (mean over channels)
- G = f_ch ⊗ f_sp ∈ ℝ^(B×C×N) [outer product]

Frobenius Normalization & Gating

- $\tilde{G} = G / (\|G\|_F + \varepsilon)$ [Frobenius norm stabilization]
- $g = \sigma(W_G \cdot \tilde{G}) \in \mathbb{R}^{(B \times C_{mid})}$ [linear projection]
- $x_{scfg} = x \odot g$ [spatial modulation]
-
- Only 82 parameters (0.03% of network)
- Captures bilinear correlations missed by first-order GAP



Minority-Aware Hybrid Loss Function

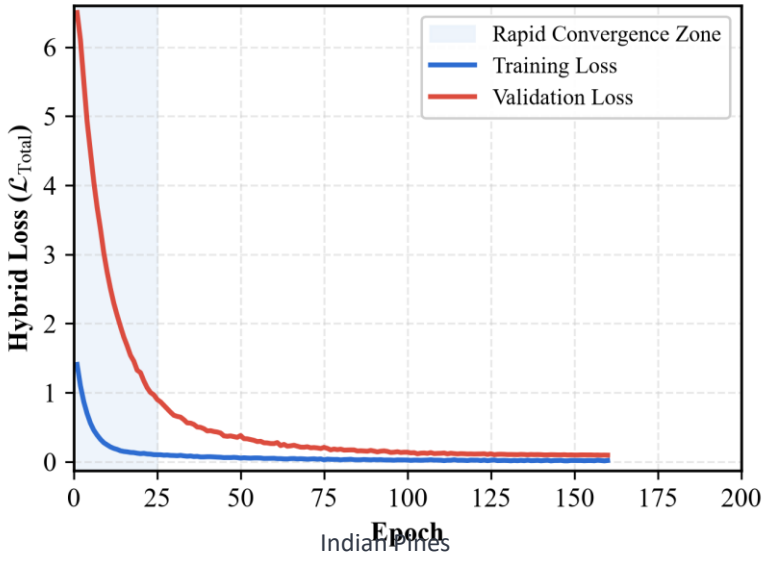
Spread Loss (weight = 0.7)

- Dynamic margin m(r) = 0.2 + 0.7·r, scheduled over training
- Early epochs: relaxed margin (m=0.2) prevents over-penalization
- Late epochs: strict margin (m=0.9) enforces wide class boundaries

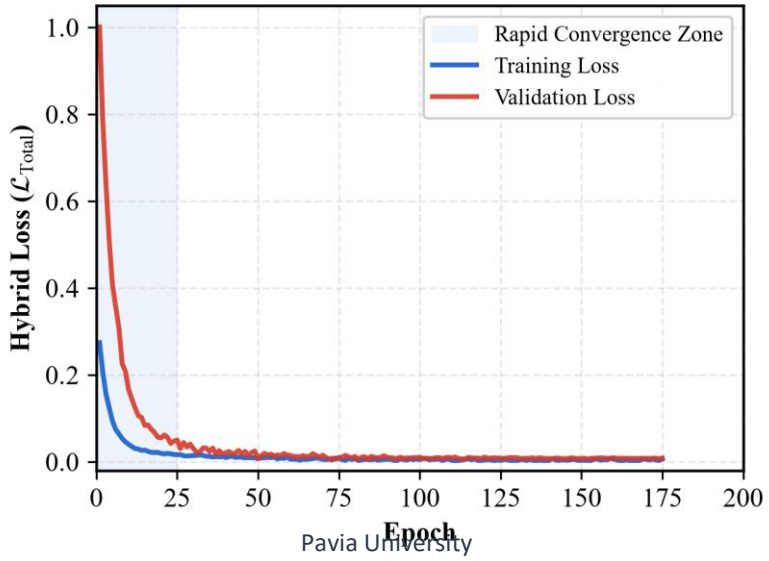
Focal Loss (weight = 0.3)

- Inverse-frequency class weights: $w_c = 1/(N_c + \varepsilon)$
- Focusing parameter γ=2.0 down-weights easy-to-classify samples
- Prevents majority classes from dominating gradients over minorities

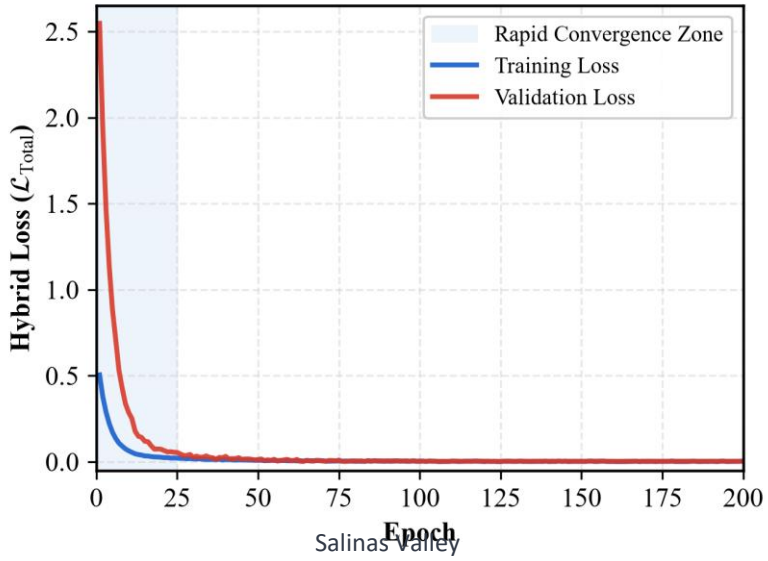
Training & Validation Loss Convergence (200 Epochs)



Indian Pines



Pavia University



Salinas Valley

Experimental Setup & Benchmark Datasets

	 Indian Pines (IN)	 Pavia University	 Salinas Scene
Sensor	AVIRIS	ROSIS	AVIRIS
Resolution	20 m/pixel	1.3 m/pixel	3.7 m/pixel
Size	145 × 145	610 × 340	512 × 217
Bands	200 (filtered)	103 channels	204 (filtered)
Classes	16	9	16
Labeled Samples	10,249	42,776	54,129

Standardized Evaluation Protocol

- Split: 10% Train / 10% Validation / 80% Test
- Stratified per class (preserved proportions)
- Patch size: 11×11 spatial neighborhood
- Batch size: 32, Epochs: 200
- Optimizer: Adam, LR: $5\times10^{-4} \rightarrow 10^{-6}$
- Cosine Annealing LR schedule
- Hardware: NVIDIA Tesla T4 GPU (Kaggle)

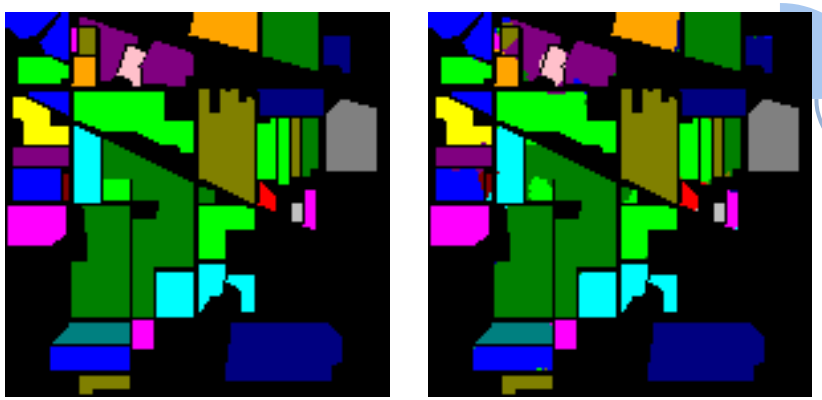
8 Baseline Models Benchmarked

- 3D CNNs: SSRN, A2S2K-ResNet
- Transformers: ViT, HSI-SSFTT, HyperRoPE-HiT
- Attention Networks: SSAN
- Capsule / Hybrid: DMC, CRV-MHDL
- All trained from scratch, identical protocol
- 24 total baseline runs (8 models × 3 datasets)

Results: Indian Pines — Per-Model Comparison

Model	OA (%)	AA (%)	κ	Time (min)
SSRN	**98.52**	79.99	0.9831	6.29
A2S2K-ResNet	98.11	94.41	0.9785	14.89
HyperRoPE-HiT	94.07	89.37	0.9324	1.43
HSI-SSFTT	98.45	94.29	0.9823	1.05
DMC	82.55	64.11	0.7989	17.73
ViT	90.33	89.12	0.8895	1.68
CRV-MHDL	67.20	51.71	0.6222	1.33
SSAN	64.19	72.77	0.5979	14.00
SGSCF-Net (Ours)	**98.05**	**94.82**	**0.9778**	**8.80**

Classification Maps



Ground Truth | SGSCF-Net

Key Finding

- SSRN tops OA (98.52%) but 0% on
- Alfalfa, Grass-mowed, Oats
- SGSCF-Net: BEST AA (94.82%)
- Alfalfa: 95.00%, Oats: 54.17%
- vs. SSRN: 0%, 0%, 0% on those

Results: Pavia University — Urban Material Classification

99.82%	99.71%	99.76	12.9 min
Overall Accuracy	Average Accuracy	Kappa (κ × 100)	Training Time



Ground Truth



SGSCF-Net

Per-Class Highlights

- 4 classes at 100.00%: Asphalt, Painted Metal, Bare Soil, Bitumen
- Meadows: 100.00% (14,919 test pixels classified perfectly)
- Gravel: 99.76% — historically confused with Asphalt
- Shadows (947 px): 99.34% via SCFG covariance operator

Baseline Context

- SSRN: 99.87% OA (marginally higher, 11.36 min)
- A2S2K-ResNet: 99.83% OA, 41.12 min (3.2× slower)
- SSAN: 76.82% OA — Gravel drops to 17.96%

Results: Salinas Valley — Fine-Grained Agriculture

99.92%	99.92%	99.91	57.7 min
Overall Accuracy	Average Accuracy	Kappa (κ × 100)	Training Time



Ground Truth



SGSCF-Net

Per-Class Highlights

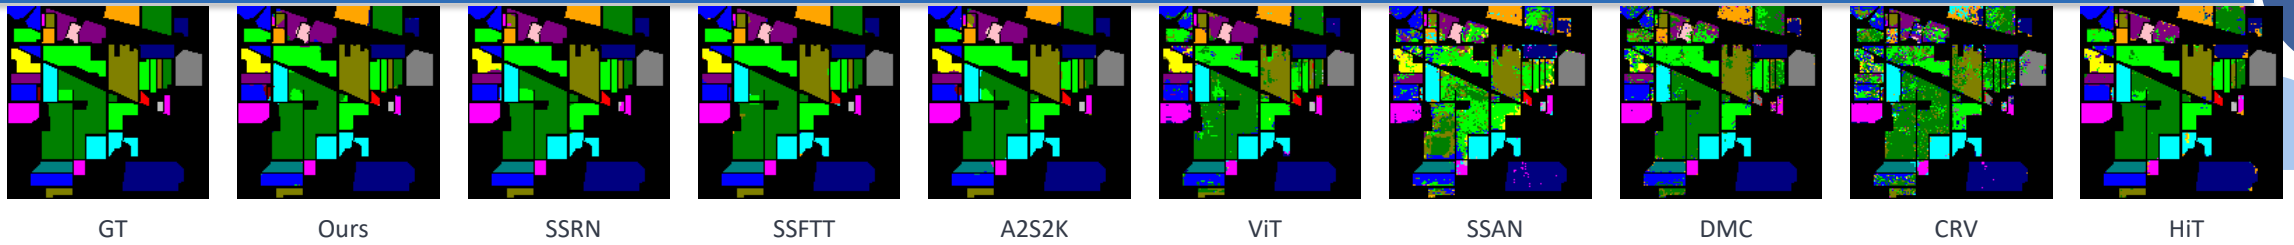
- 11 of 16 classes reach 100.00% accuracy
- All 4 Romaine lettuce growth stages (4-7 wk): 99.65-100%
- Vineyard untrained: 99.73% vs. Vineyard trellis: 100.00%
- SCFG captures chlorophyll absorption band-ratio shifts

Baseline Context

- A2S2K-ResNet: 100.00% OA (10.34 min, similar speed)
- SSAN: 88.20% OA — Grapes untrained 57.89%, collapse
- OA trend: 98.05% (IP) → 99.82% (PU) → 99.92% (SV)

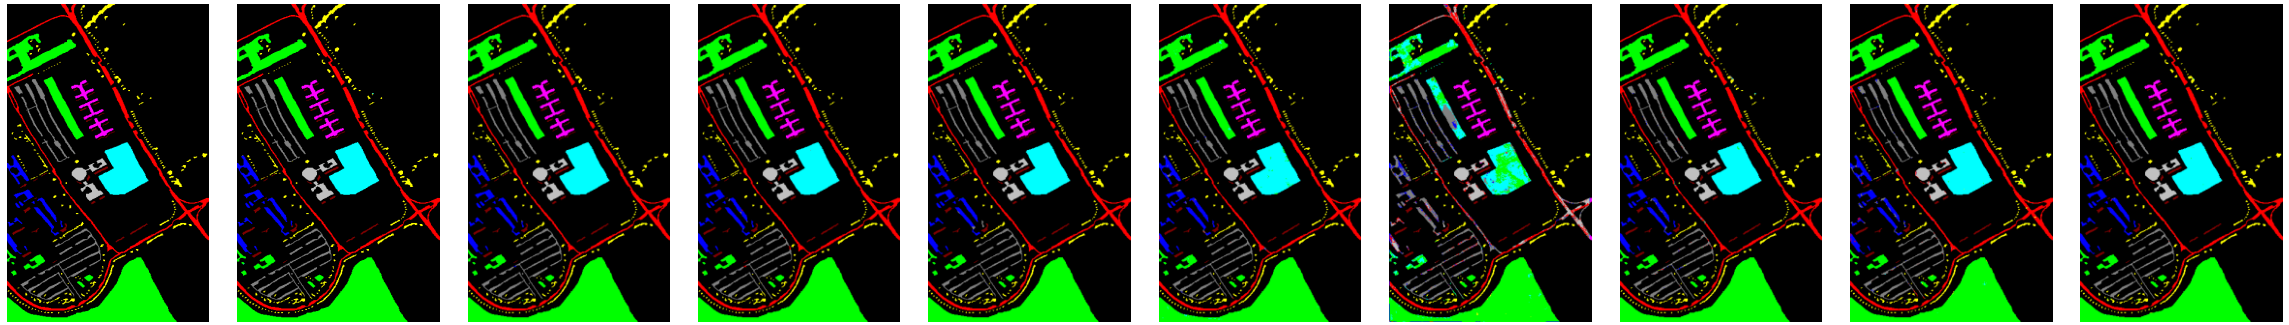
Visual Comparison: Classification Maps Across Models

Indian Pines



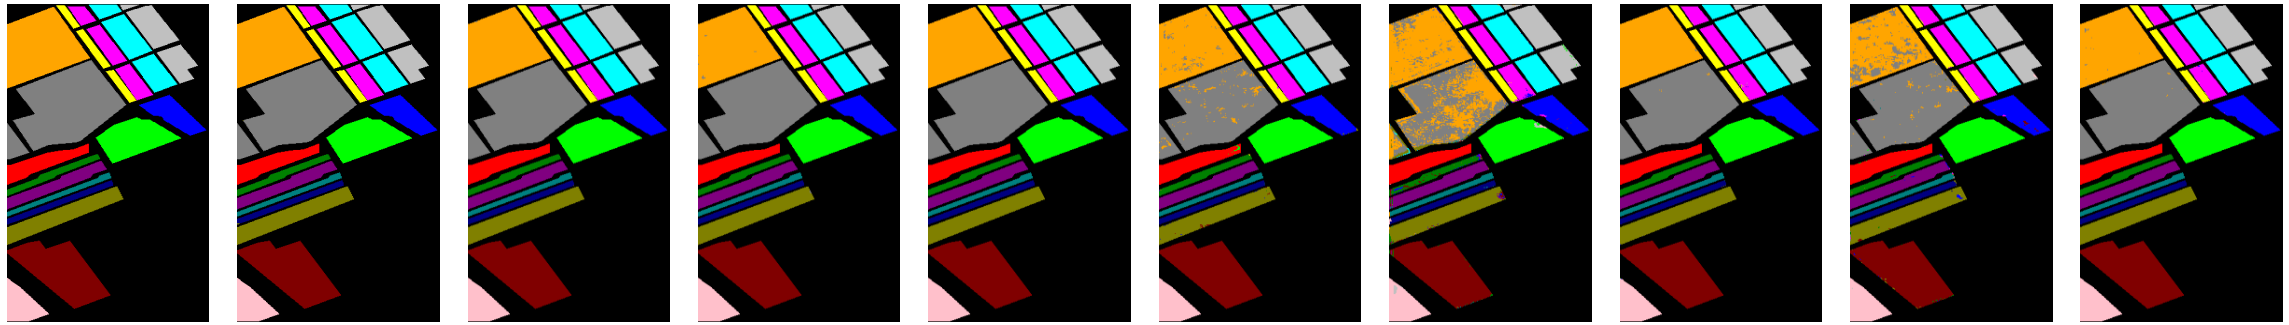
GT | Ours | SSRN | SSFTT | A2S2K | ViT | SSAN | DMC | CRV | HiT

Pavia Univ.



GT | Ours | SSRN | SSFTT | A2S2K | ViT | SSAN | DMC | CRV | HiT

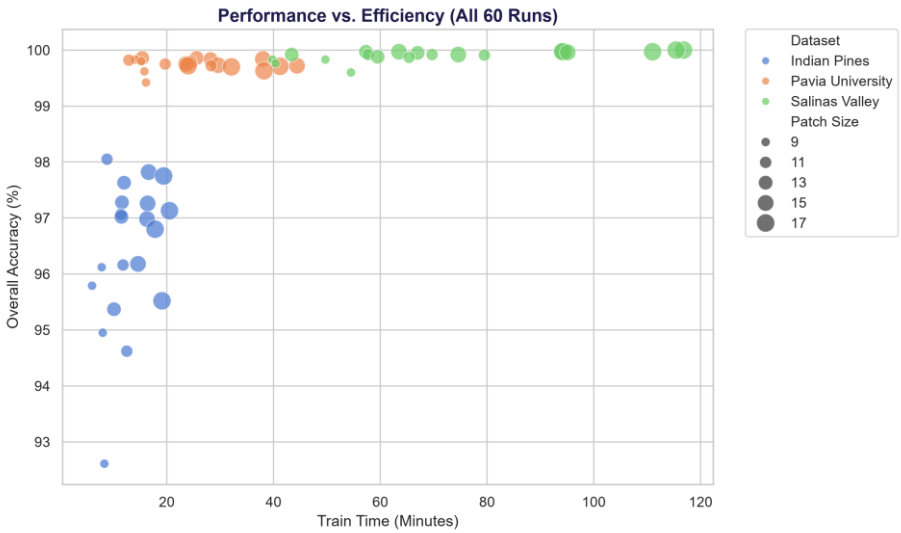
Salinas



GT | Ours | SSRN | SSFTT | A2S2K | ViT | SSAN | DMC | CRV | HiT

Computational Efficiency & Parameter Analysis

Model	Params	Checkpoint	IP Time	OA (%)
ViT	~610K	2.4 MB	1.68 min	90.33
HSI-SSFTT	~156K	625 KB	1.05 min	98.45
A2S2K-ResNet	~380K	1.5 MB	14.89 min	98.11
SSAN	6.40M	—	14.00 min	64.19
DMC	~480K	—	17.73 min	82.55
SGSCF-Net (Ours)	**308,097**	**1.2 MB**	**8.80 min**	**98.05**

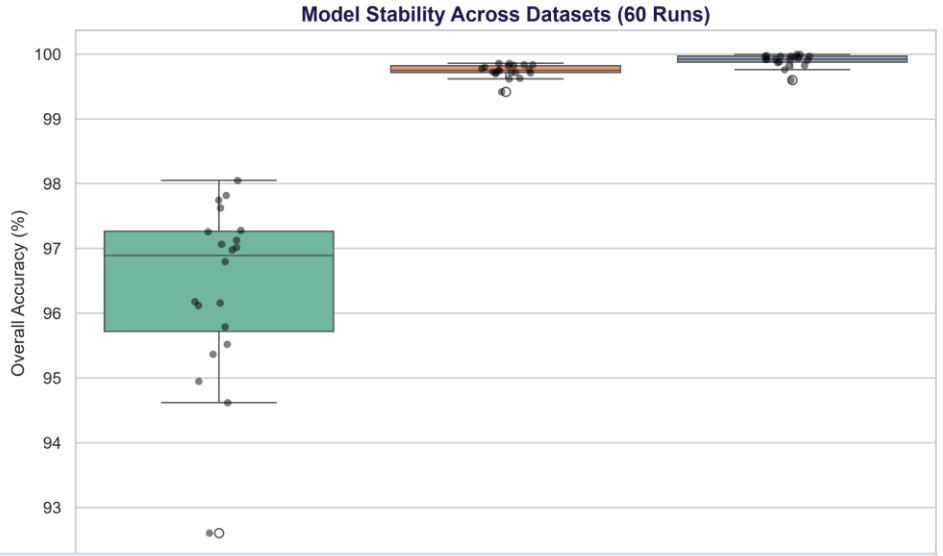


Module-Level Parameter Breakdown (Indian Pines)

- SBAM: 10,225 (3.32%) | CBSA: 80 (0.03%) | SCFG: 82 (0.03%) → 3 novel modules = only 10,387 params (3.37% of total)
- Primary volume: Spatial 3D entry (11.98%) + Spatial ResBlocks (81.00%) — standard Conv3D backbone
- Classifier head: 272 params (0.09%) — single linear layer, no deep FC stacks

Ablation Study & Block Depth Sensitivity

#	SBAM	CBSA	SCFG	H.Loss	OA (%)	AA (%)
1	✓	×	×	×	93.65	89.24
2	✓	✓	×	×	95.18	91.80
3	✓	✓	✓	×	96.44	93.05
4	✓	✓	✓	CE	97.26	93.78
5	**✓**	**✓**	**✓**	**✓**	**98.05**	**94.82**



Ablation Insights

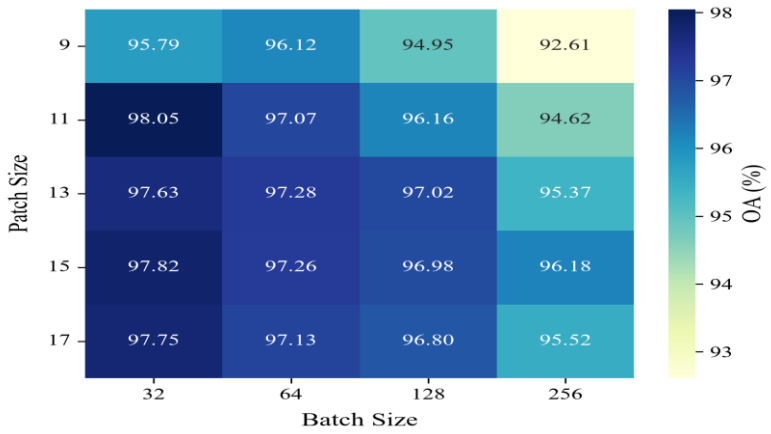
- SBAM alone lifts baseline from ~90% to 93.65% OA
- Replacing shared SE gate with CBSA: +1.53% OA
- SCFG adds +1.26% via second-order correlations
- Hybrid loss: +0.79% OA, +1.04% AA (minority protection)
- Oats class: 0% (CE-only) → 54.17% (hybrid loss)

Block Depth Sensitivity (D)

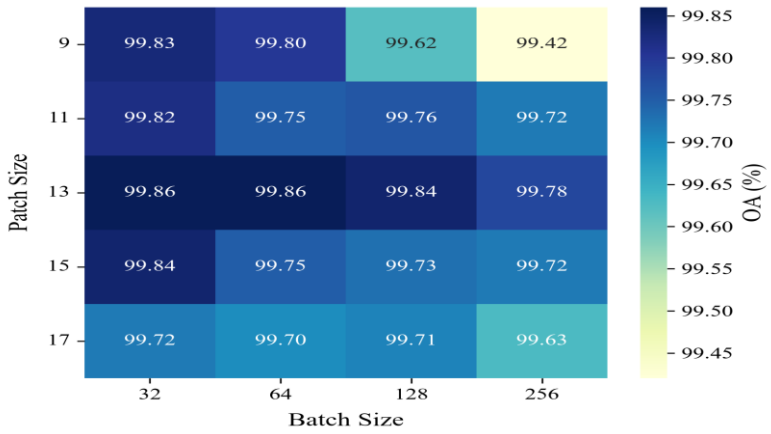
- D=1: 179K params, 95.82% OA (underfits)
- D=2: 308K params, 98.05% OA (optimal ✓)
- D=3: 437K params, 97.88% OA (diminishing returns)
- D=4: 566K params, 97.40% OA (overfitting onset)
- D=2 yields best accuracy-efficiency trade-off

60-Run Hyperparameter Grid Search

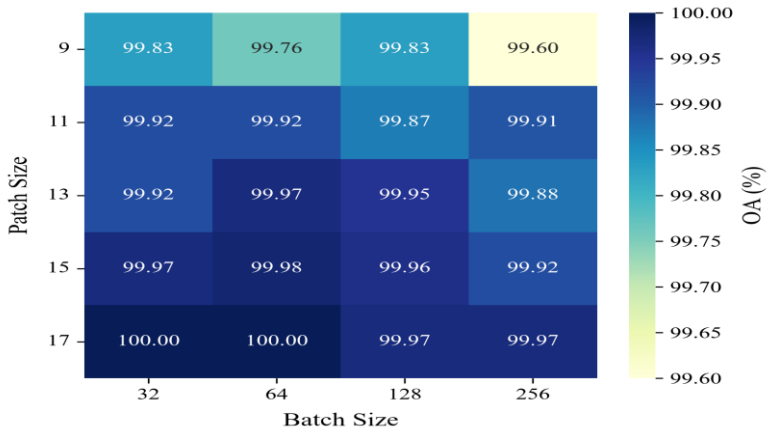
5 Patch Sizes (9×9, 11×11, 13×13, 15×15, 17×17) × 4 Batch Sizes (32, 64, 128, 256) = 20 runs/dataset × 3 datasets = 60 total



Indian Pines OA Heatmap



Pavia University OA Heatmap



Salinas Valley OA Heatmap

Optimal Configuration

- Patch 11×11, Batch 32
- OA: 98.05% (IP), 99.82% (PU), 99.92% (SV)
- Consistent Rank 1 or within 0.08%

Spatial Window Trade-off

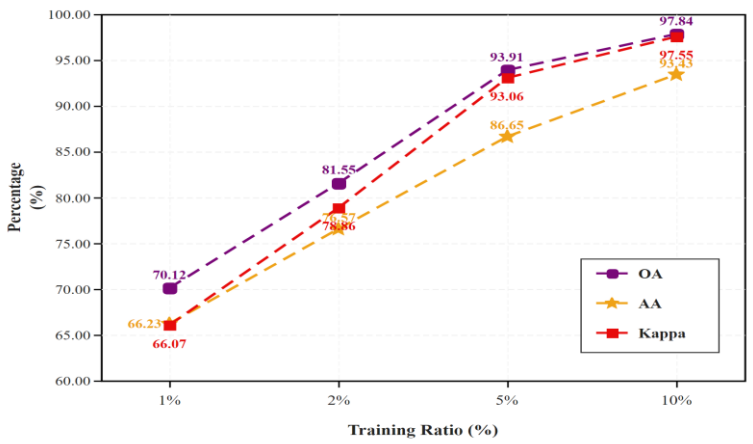
- 9×9: underfits (misses texture boundaries)
- 15×15, 17×17: overfits (cross-field contamination)
- 11×11 at 20m/pixel = 220m ground distance

Batch Size Dynamics

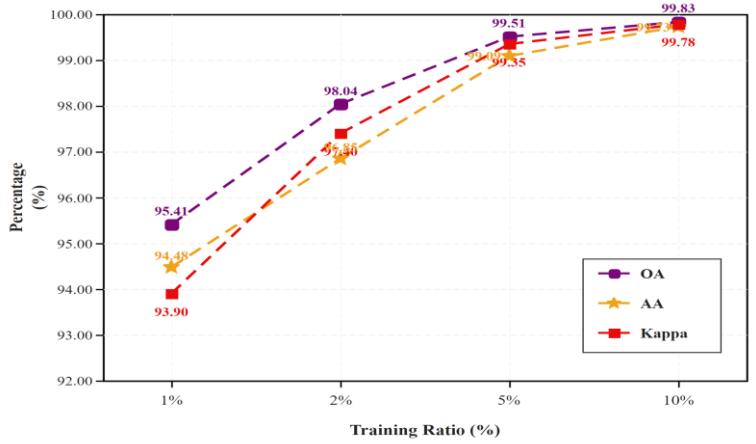
- B=32 beats B=256 by 3.43% on Indian Pines
- Small batches: higher gradient variance
- Preserves minority class decision boundaries

Data Scarcity Stress Testing: 1% – 10% Training Ratios

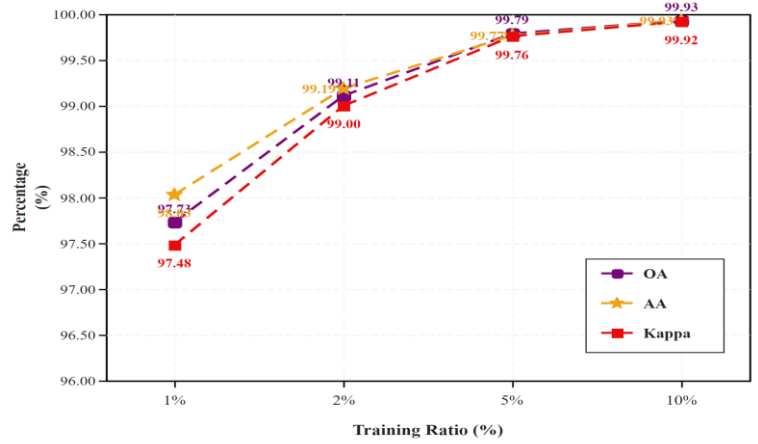
Dataset	1% Train	2% Train	5% Train	10% Train	Drop
Indian Pines	70.12%	81.55%	93.91%	97.84%	-27.73%
Pavia University	95.41%	98.04%	99.51%	99.83%	-4.43%
Salinas Valley	97.73%	99.11%	99.79%	99.93%	-2.20%



Indian Pines Degradation



Pavia University Degradation



Salinas Valley Degradation

Salinas: most resilient (−2.20%, 97.73% at 1%) | Pavia: moderate (−4.43%, 95.41% at 1%) | Indian Pines: steepest drop (−27.73%, 70.12% at 1% — Oats has 1–2 patches)

Cross-Dataset Summary & Key Achievements

98.05%	99.82%	99.92%	308K
Indian Pines OA	Pavia Univ. OA	Salinas Valley OA	Total Parameters

Project Achievements

- Designed & implemented 3 novel modules (SBAM, CBSA, SCFG)
- Best Average Accuracy on Indian Pines (94.82%) across all 9 models
- 99.82% OA on Pavia, 99.92% OA on Salinas (near-perfect)
- 120+ GPU runs: 60 grid search + 12 scarcity + 24 baseline + proposed
- O(N) linear complexity — no quadratic self-attention overhead
- 1.18s test inference on Indian Pines (vs. 135.74s SSAN)
- Complete 31-page ACM-format research manuscript drafted

Why SGSCF-Net Matters

- Lightweight: 308K params vs. 6.4M SSAN, ~610K ViT, ~480K DMC
- Balanced: Protects minority classes where SSRN/SSAN collapse to 0%
- Generalizable: Consistent top-tier across 3 diverse HSI scenes
- Efficient: 8.80 min training vs. 17.73 (DMC), 41.12 (A2S2K)
- Modular: SBAM/CBSA/SCFG are plug-and-play for any 3D CNN
- Reproducible: Fixed 10/10/80 split, 200 epochs, public datasets
- Edge-deployable: 1.2 MB checkpoint, suitable for UAV/satellite

Future Work & Research Directions

1 Hardware-Aware Quantization

INT8/FP8 tensor quantization + structured pruning for real-time inference on edge NPUs and FPGAs aboard UAVs and satellite platforms

2 Self-Supervised Pretraining

Masked spectral pre-training on large uncurated HSI cubes to learn transferable representations and reduce label dependency below 1%

3 Multi-Modal Fusion

Extend SCFG covariance mechanism to joint HSI + LiDAR classification for multi-sensor Earth observation pipelines

4 Journal Publication

Target IEEE Transactions on Geoscience and Remote Sensing (TGRS) or similar Q1 journal for peer-reviewed publication of this work

Thank You

Questions & Discussion

SGSCF-Net: Spectral-Guided Spatial Cross-Attention Fusion Network for Hyperspectral Image Classification

Presented By

Tejas Surya — Roll No. 23118098

B. Vyshnavi — Roll No. 23118021

Valluri Lahitha — Roll No. 23118106

Guided By

Dr. R.R. Janghel

Associate Professor, Dept. of IT

National Institute of Technology, Raipur

National Institute of Technology, Raipur • Department of Information Technology • November 2026